

randomForest

Random Forest Tests

Here I am playing around with Random Forest Classification in a data set I make up.

The data set

First, I am setting up the session and the data set. I am essentially imagining a sample of individuals, for whom we know age, weight, an affected trait, maybe fitness (V02max), and have genotyped (0/1/2) four loci. We also have the phenotype of interest, perhaps a disease (0/1) that affects fitness (age and weight do, too), and is in turn caused by three of the four genotyped loci. I wanted to see how a random forest would deal with this scenario.

```
library(tidyverse)
```

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr      1.1.1      v readr      2.1.4
v forcats    1.0.0      v stringr    1.5.0
v ggplot2    3.4.1      v tibble     3.2.1
v lubridate  1.9.2      v tidyr      1.3.0
v purrr      1.0.1
-- Conflicts ----- tidyverse_conflicts() --
x dplyr::filter() masks stats::filter()
x dplyr::lag()     masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become
```

```
library(randomForest)
```

randomForest 4.7-1.1

Type `rfNews()` to see new features/changes/bug fixes.

Attaching package: 'randomForest'

The following object is masked from 'package:dplyr':

`combine`

The following object is masked from 'package:ggplot2':

`margin`

```
count_individuals = 2000
count_loci = 4

genotype_matrix <-
  matrix(
    rbinom(count_individuals * count_loci, 1, 0.5) + rbinom(count_individuals * count_loci,
    nrow = count_individuals,
    ncol = count_loci
  )
genotype_df <- as_tibble(genotype_matrix)
```

Warning: The `x` argument of `as_tibble.matrix()` must have unique column names if
`.name_repair` is omitted as of tibble 2.0.0.

i Using compatibility `.name_repair`.

```
colnames(genotype_df) <- paste0("locus_", 1:count_loci)

additional_data <-
  tibble(
    ID = 1:count_individuals,
    weight = rnorm(count_individuals, mean = 80, sd = 5),
    age = sample(20:80, count_individuals, replace = TRUE)
  )

data_combined <- bind_cols(additional_data, genotype_df)
```

```

make_binary_trait = function(row) {
  # here are the rules how the binary trait is generated
  data_combined_slice <- slice(data_combined, row)
  #print(data_combined_slice)
  individual_risk <- 0
  individual_risk <-
    individual_risk + (pull(data_combined_slice, locus_1)) * 0.1 # additive effect
  individual_risk <-
    individual_risk + (pull(data_combined_slice, locus_2)) * 0.1
  individual_risk <-
    individual_risk + (pull(data_combined_slice, locus_3)) * 0.1
  individual_risk <-
    individual_risk + ((pull(data_combined_slice, locus_1) == 2) * (pull(data_combined_sli
                                                                    2)) * 0.4

  individual_risk <-
    individual_risk + ((pull(data_combined_slice, locus_1) == 2) * (pull(data_combined_sli
                                                                    2) * (pull(data_comb

  if (individual_risk > 1) {
    individual_risk <- 1
  }
  phenotype <- rbinom(1, 1, individual_risk)
  return(phenotype)
}

add_affected_trait <- function(row) {
  data_combined_slice <- slice(data_combined, row)
  mean <-
    40 # the average in the healthy population of a trait, say V02max
  mean <-
    mean - (pull(data_combined_slice, phenotype) == 1) * 5 # decrease in 5 for those with
  mean <- mean - pull(data_combined_slice, weight) * 0.1 # decrease with weight
  mean <- mean - pull(data_combined_slice, age) * 0.1 # decrease with age
  return(rnorm(1, mean, sd = 3))
}

# add the phenotype and affected trait to the data frame

data_combined$phenotype <-
  as.factor(unlist(lapply(
    1:nrow(data_combined), make_binary_trait
  )))

```

```

data_combined$affected_trait <-
  unlist(lapply(1:nrow(data_combined), add_affected_trait))
data_combined$weight = as.numeric(scale(data_combined$weight, center = TRUE, scale =
                                     TRUE))
data_combined$age = as.numeric(scale(data_combined$age, center = TRUE, scale =
                                     TRUE))
data_combined$affected_trait = as.numeric(scale(
  data_combined$affected_trait,
  center = TRUE,
  scale = TRUE
))

summary(data_combined$phenotype)

```

```

      0      1
1398  602

```

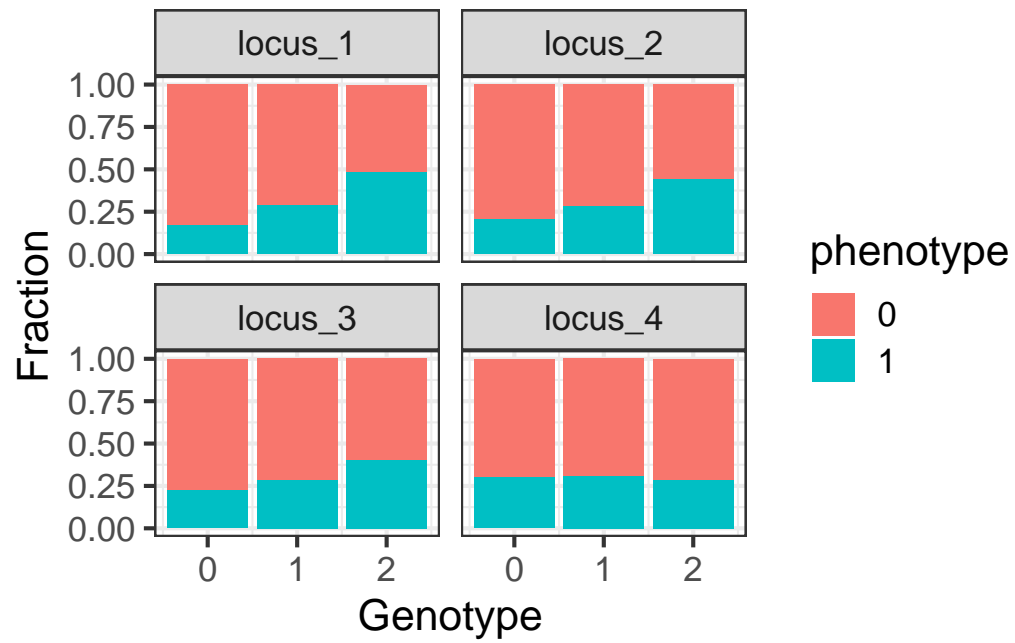
Above, you can see the distribution of the phenotype. A bit imbalanced, but still plenty of samples.

Let's have a look at the traits we have. **But note that they have been scaled for better modelling.**

```


library(ggplot2)
ggplot(data_combined%>%pivot_longer(cols=c(locus_1,locus_2,locus_3,locus_4)), aes(value, g
  geom_bar(position="fill")+facet_wrap(~name)+xlab("Genotype")+ylab("Fraction")+theme_bw

```



```
ggplot(data_combined, aes(age, fill = phenotype)) +
  geom_histogram(position = "fill")+theme_bw(16)+ylab("Fraction")
```


`stat_bin()` using `bins = 30`. Pick better value with `binwidth`.



```
ggplot(data_combined, aes(weight, fill = phenotype)) +
  geom_histogram(position = "fill")+theme_bw(16)+ylab("Fraction")
```

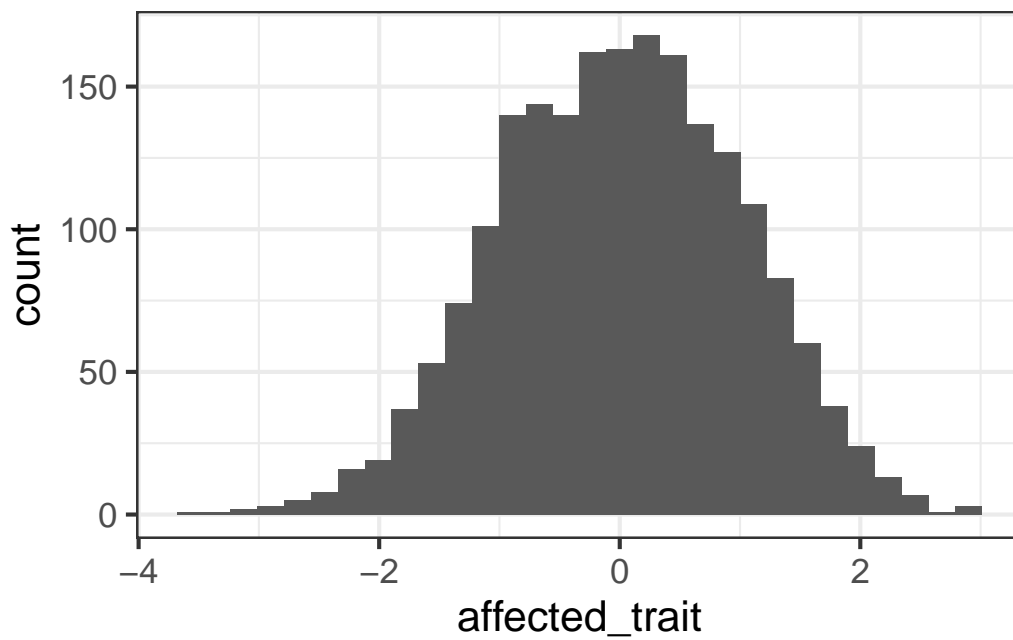
``stat_bin()`` using ``bins = 30``. Pick better value with ``binwidth``.

Warning: Removed 2 rows containing missing values (``geom_bar()``).



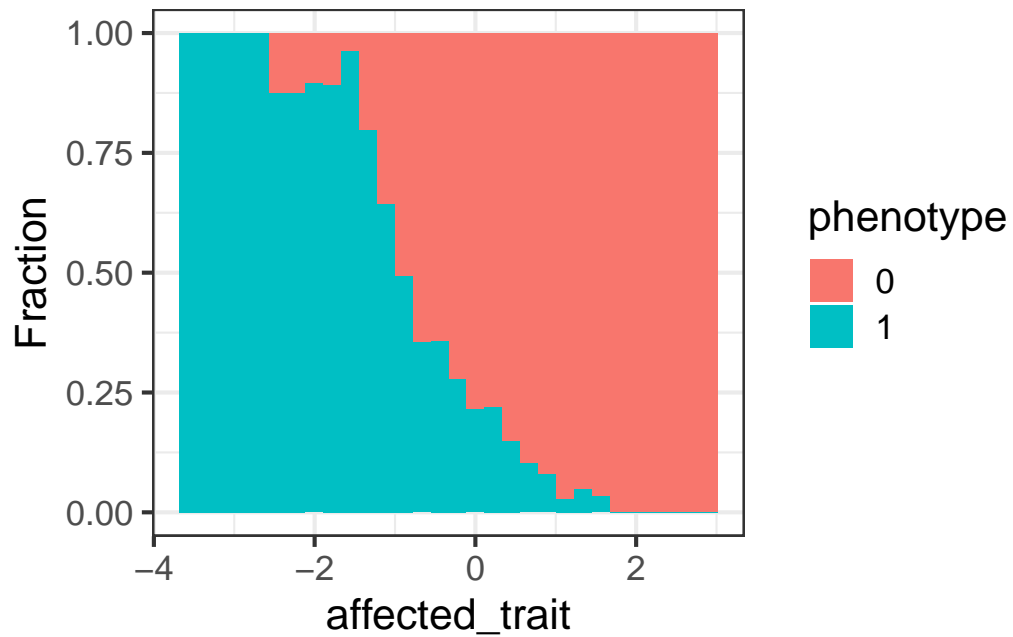
```
ggplot(data_combined, aes(affected_trait)) +  
  geom_histogram()+theme_bw(16)
```

`stat_bin()` using `bins = 30`. Pick better value with `binwidth`.




```
ggplot(data_combined, aes(affected_trait, fill = phenotype)) +  
  geom_histogram(position = "fill")+theme_bw(16)+ylab("Fraction")
```


`stat_bin()` using `bins = 30`. Pick better value with `binwidth`.



```
ggplot(data_combined, aes(weight, affected_trait, color = phenotype)) +  
  geom_point(shape = 1)+theme_bw(16)
```

```
ggplot(data_combined, aes(age, affected_trait, color = phenotype)) +  
  geom_point(shape = 1)+theme_bw(16)
```




OK, definitely realistic looking and tons of noise!

I will first see if an unsupervised forest finds some relevant structure in the data.

```
data_combined.urf <-  
  randomForest(select(data_combined,-1,-(ncol(data_combined) - 1)), ntree =  
               count_individuals * 10)  
MDSplot(data_combined.urf, data_combined$phenotype)
```

Warning in RColorBrewer::brewer.pal(nlevs, "Set1"): minimal value for n is 3, returning requested



Without any supervision, we do not see strong structuring along the phenotype (color), but some can be seen.

Can we get a good prediction with a random forest classification after some tuning?

```
library(caret)
```

Loading required package: lattice

Attaching package: 'caret'

The following object is masked from 'package:purrr':

```
lift
```

```

# split the data
ids <- createDataPartition(data_combined$phenotype, p = 0.8, list = F)
train_set <- data_combined[ids, ]
test_set <- data_combined[-ids, ]

cn <- trainControl(method = "cv", number = 10)

grid <- expand.grid(mtry = 2:(ncol(train_set) - 1))

fit <-
  train(
    phenotype ~ .,
    data = train_set,

    method = "rf",
    trControl = cn,
    tuneGrid = grid,
    ntree = count_individuals * 10,
    maxnodes = 500
  )

```

Let's see how well the model works!

```

p <- predict(fit, test_set %>% select(-(ncol(test_set) - 1)))

print(cM<-confusionMatrix(p, test_set$phenotype))

```

Confusion Matrix and Statistics

	Reference	
Prediction	0	1
0	261	49
1	18	71

Accuracy : 0.8321
 95% CI : (0.7917, 0.8674)
 No Information Rate : 0.6992
 P-Value [Acc > NIR] : 7.523e-10

 Kappa : 0.569

McNemar's Test P-Value : 0.0002473

Sensitivity : 0.9355
Specificity : 0.5917
Pos Pred Value : 0.8419
Neg Pred Value : 0.7978
Prevalence : 0.6992
Detection Rate : 0.6541
Detection Prevalence : 0.7769
Balanced Accuracy : 0.7636

'Positive' Class : 0

So, concluding for now, we have built a random forest that is able to find 93.55% of individuals with an imaginary disease that decreases their physical fitness and is determined by three loci in an additive and epistatic fashion. We would probably not be able to use this model on a large number of loci because of performance, but I am curious how well it still works when there are more loci (of course, building a new forest with those first).